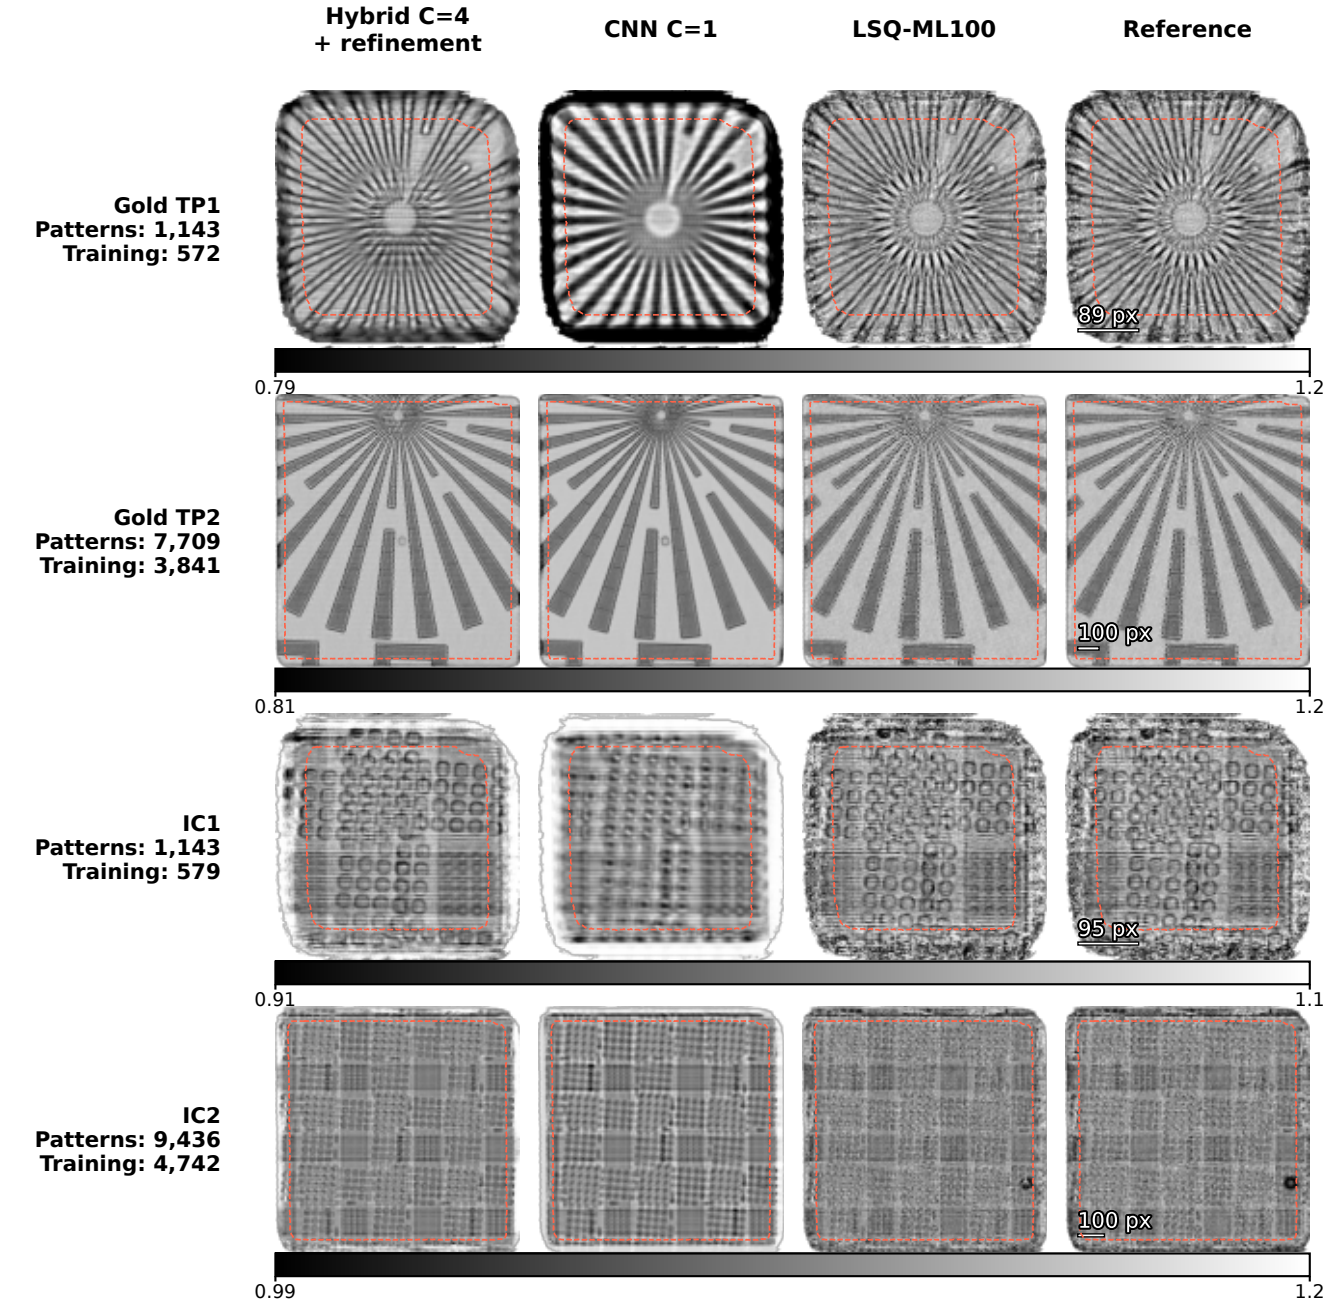
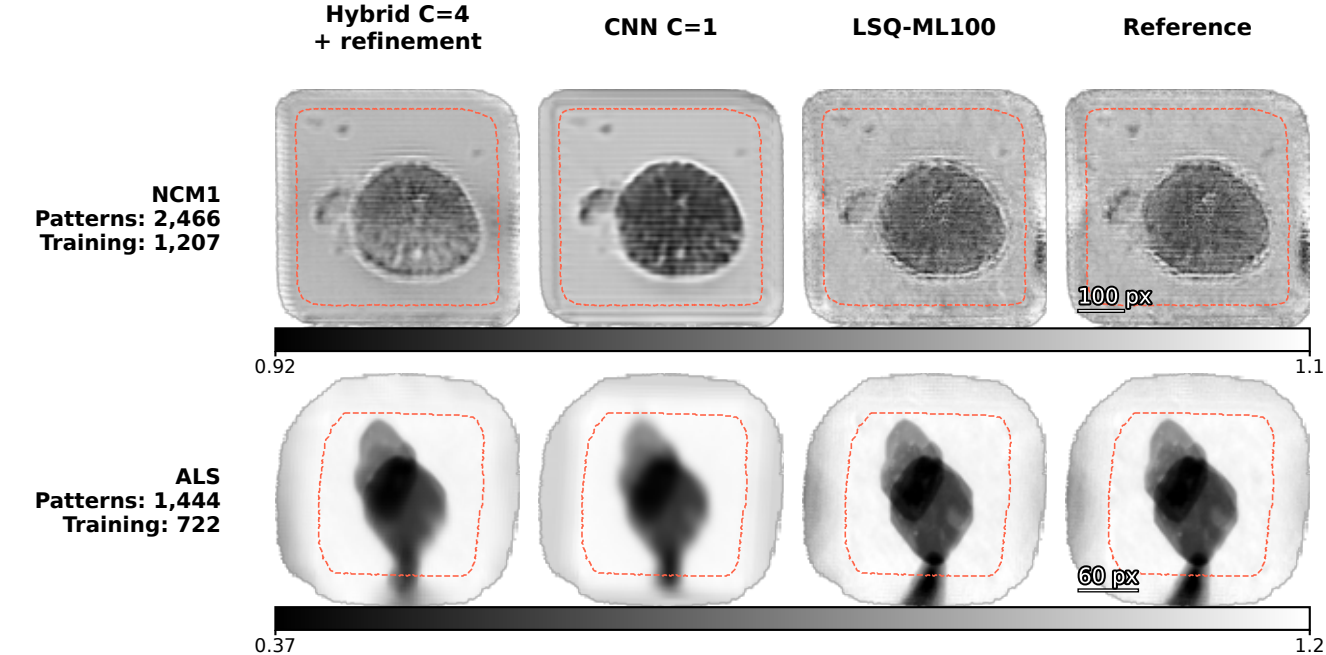


Phase is shown in radians; amplitude shows relative object amplitude. Pattern counts refer to the displayed acquisition and distinct training patterns. High half selects checkpoints (spatial validation, not an independent test; no guard band). Reference is an experimental estimate; dashed contour marks scored interior. Matched measured CNN C=1 is shown without patch-offset alignment, tilt seed, or phase-field refinement. It is an architecture baseline under this study protocol, not a reproduction of the published CNN result. Whole-acquisition images include the training region. Paired groups assign C4 only where it has lower phase MAE after refinement and higher phase SSIM; other datasets use C1. This is the C4-selected group. C4 uses 10 groups per training center; training compute is not matched to C1. TP2 is shown full field only; FRC is excluded because of scan-grid artefacts.



Phase is shown in radians; amplitude shows relative object amplitude. Pattern counts refer to the displayed acquisition and distinct training patterns. High half selects checkpoints (spatial validation, not an independent test; no guard band). Reference is an experimental estimate; dashed contour marks scored interior. Matched measured CNN C=1 is shown without patch-offset alignment, tilt seed, or phase-field refinement. It is an architecture baseline under this study protocol, not a reproduction of the published CNN result. Whole-acquisition images include the training region. Paired groups assign C4 only where it has lower phase MAE after refinement and higher phase SSIM; other datasets use C1. This is the C4-selected group. C4 uses 10 groups per training center; training compute is not matched to C1. TP2 is shown full field only; FRC is excluded because of scan-grid artefacts.