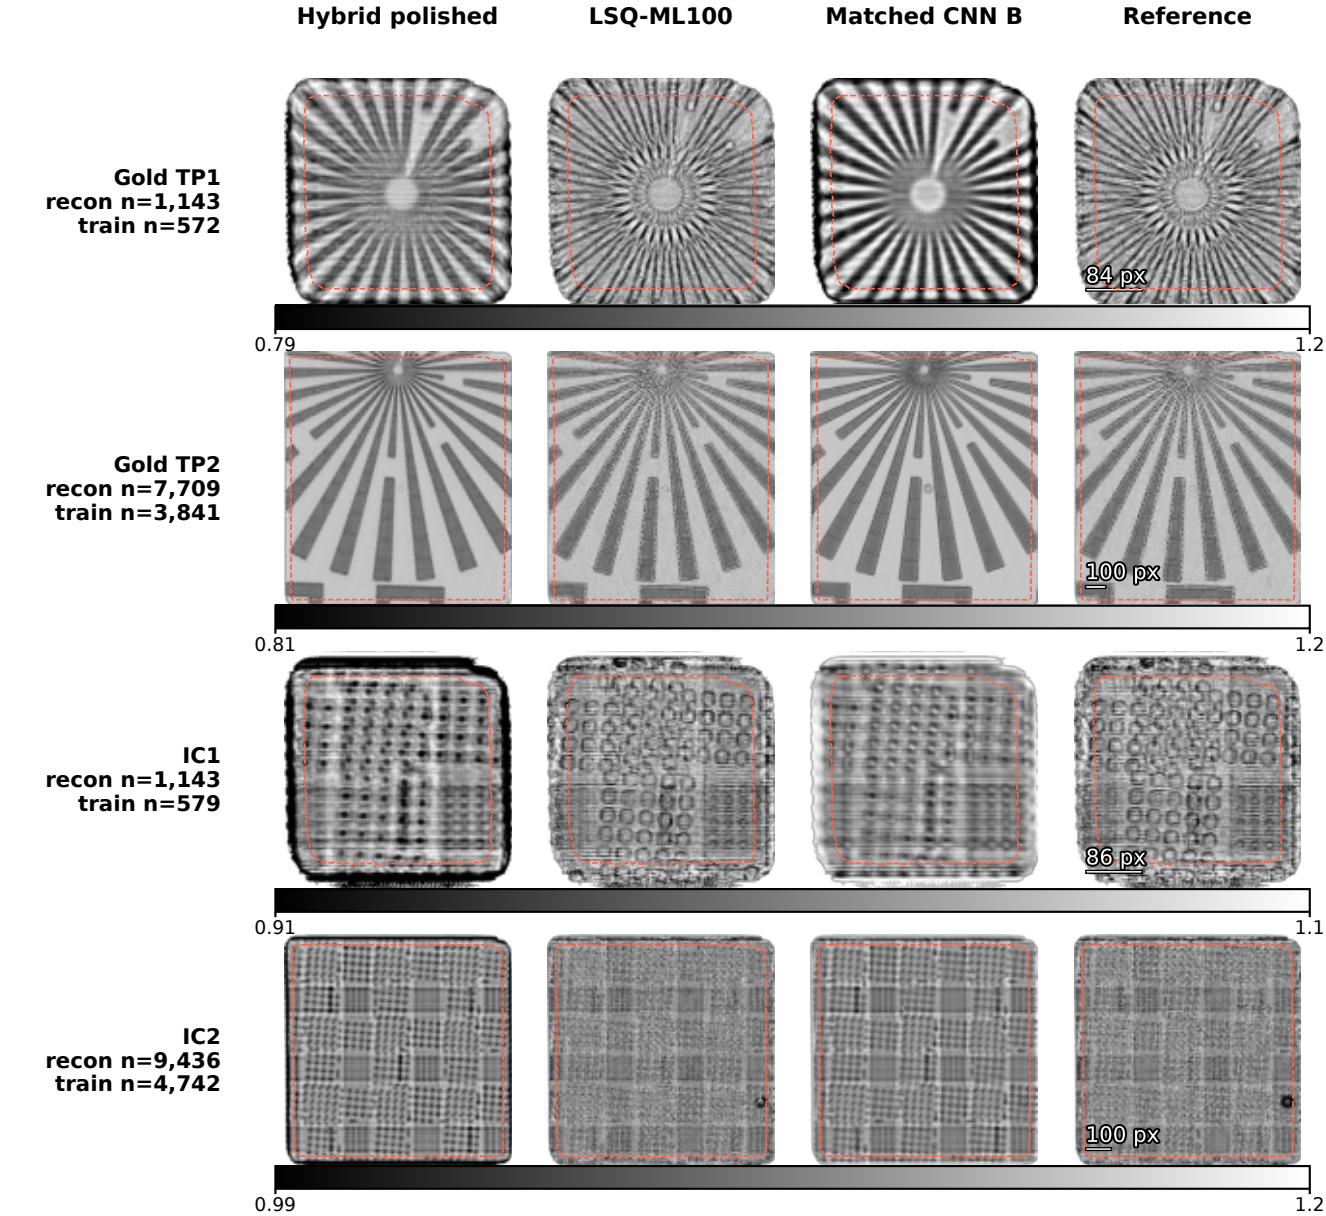
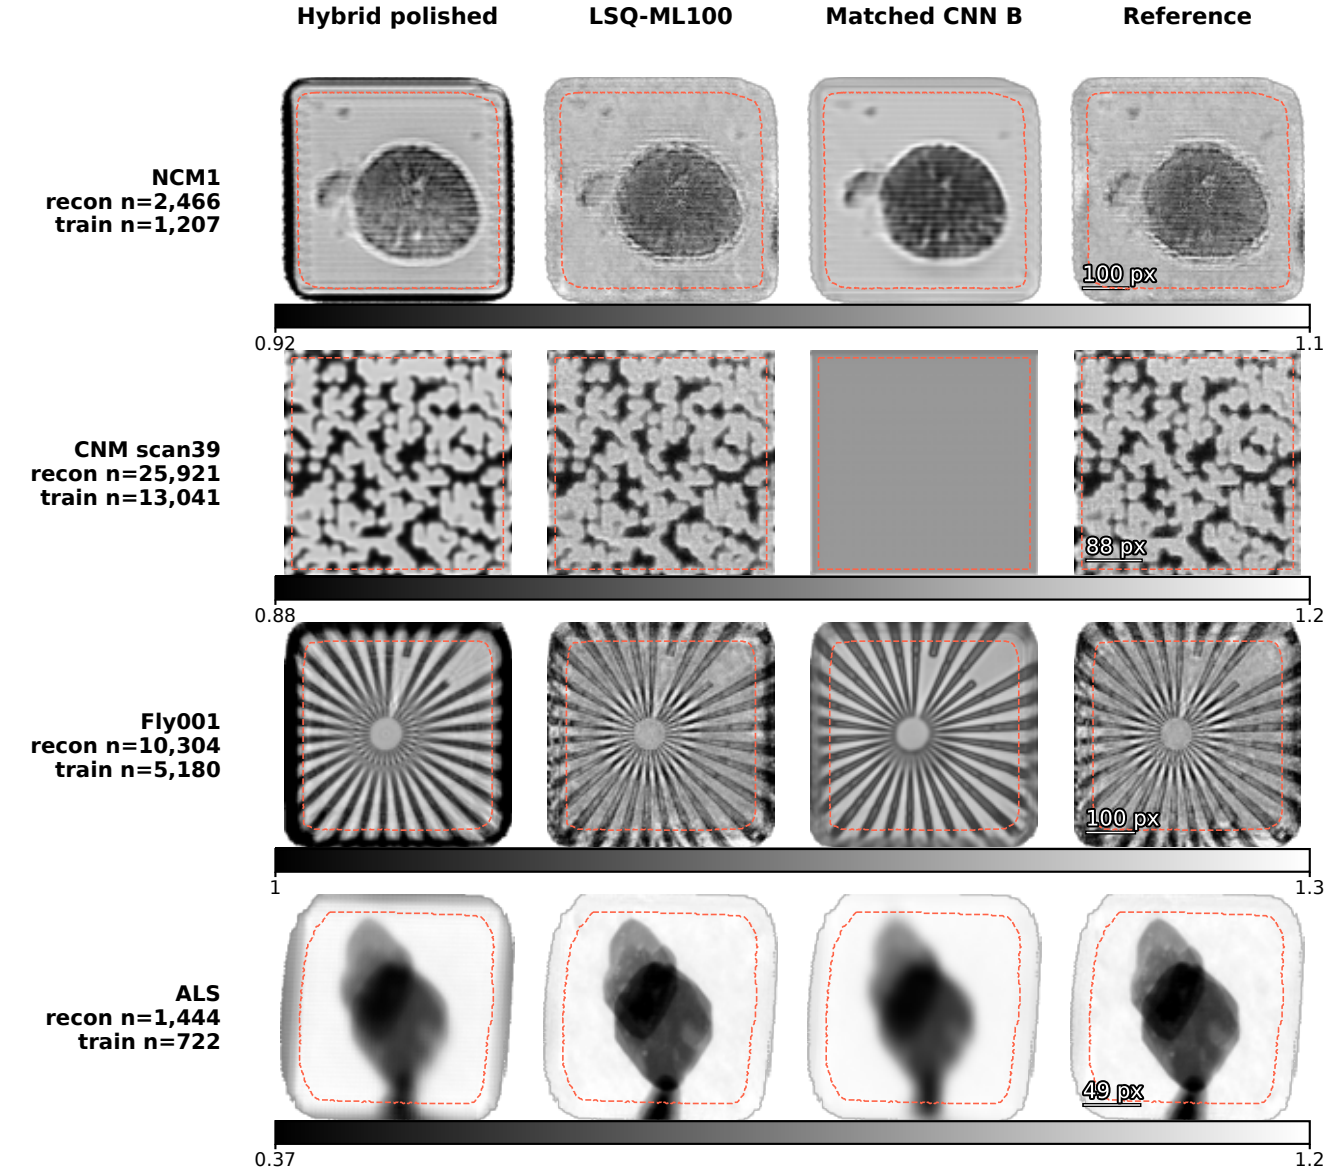
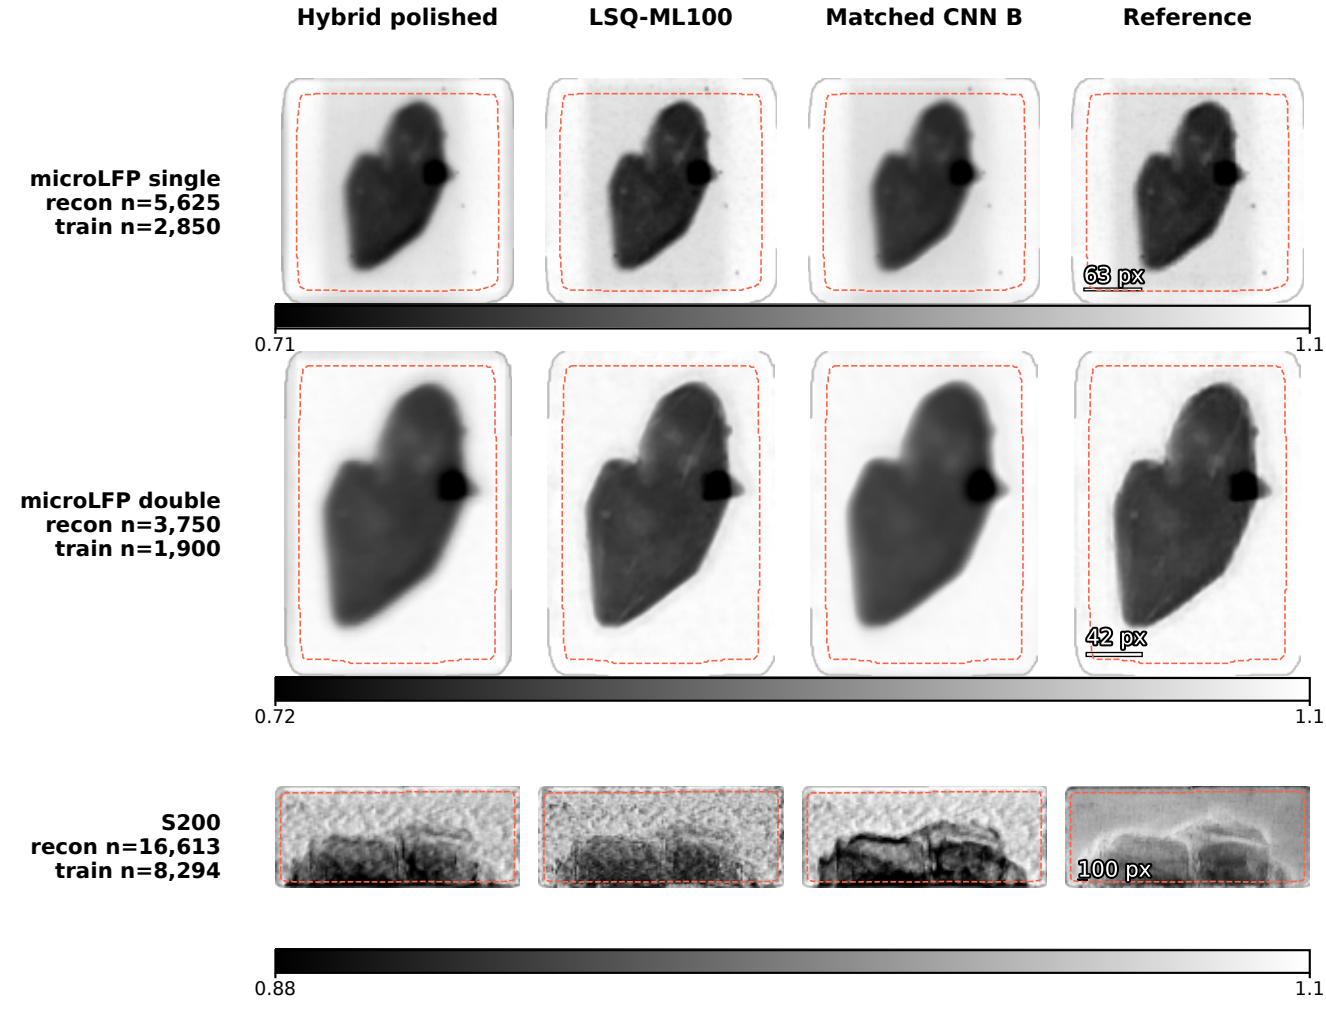


Phase is shown in radians; amplitude shows relative object amplitude. Recon n is the displayed acquisition; train n is the recorded training count. High half selects checkpoints (spatial validation, not an independent test; no guard band). Reference is an experimental estimate; dashed contour marks scored interior. Matched measured CNN C=1 is shown without patch-offset alignment, tilt seed, or phase-field refinement. It is an architecture baseline under this study protocol, not a reproduction of the published CNN result. Whole-acquisition images include the training region.



Phase is shown in radians; amplitude shows relative object amplitude. Recon n is the displayed acquisition; train n is the recorded training count. High half selects checkpoints (spatial validation, not an independent test; no guard band). Reference is an experimental estimate; dashed contour marks scored interior. Matched measured CNN C=1 is shown without patch-offset alignment, tilt seed, or phase-field refinement. It is an architecture baseline under this study protocol, not a reproduction of the published CNN result. Whole-acquisition images include the training region.



Phase is shown in radians; amplitude shows relative object amplitude. Recon n is the displayed acquisition; train n is the recorded training count. High half selects checkpoints (spatial validation, not an independent test; no guard band). Reference is an experimental estimate; dashed contour marks scored interior. Matched measured CNN C=1 is shown without patch-offset alignment, tilt seed, or phase-field refinement. It is an architecture baseline under this study protocol, not a reproduction of the published CNN result. Whole-acquisition images include the training region.